

# Pushy Postgres and Python

## BrightLink

Help people prove themselves worthy

↓

Candidate/credential &

Performance exam management

- small business
- candidate & credential management
- Python stack

Start with

## Why?

You

- *always on, social* users
- scale out: throwing hardware at the problem only goes so far
- more visibility w/o hurting performance

Me

- love Python
- love sharing knowledge

## tl;dr

Use Postgres?

Get data-driven push notices in your distributed app for free!

🐘,

@drocco007

- This is the one-slide version of this talk
- add distrib. message processing with
- data-driven push notices
- no additional infrastructure, ops changes, etc.

Outline:

- overview of Observer
- Notify/Listen in PG
- Case study

## Demand DB

- Do username & password match?
- How many people registered for event?
- Find email addresses for all users in OH

## Ubiquitous Connectivity

“Don’t call us, we’ll call you.”

New paradigm

## SQL?

Just imagine...

```
while True:
    db.execute('SELECT are_we '
               'FROM there_yet;')
    sleep(10)
```


(if all the apps on your phone did this...)

## You are here



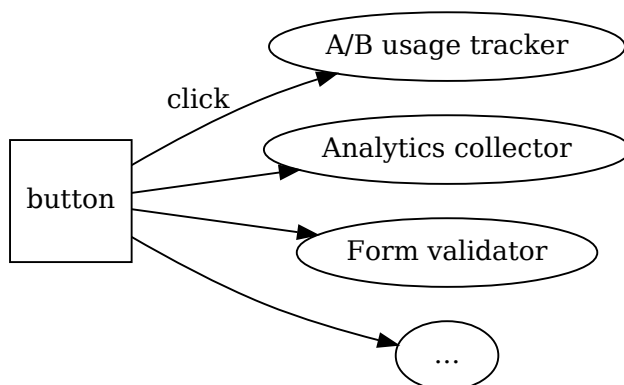
responsiveness  $\Downarrow$  system health

## Are you watching closely?



- Observer pattern
- Subject: thing being observed
- Subject: doesn't need to know about observer purpose, #
- Encapsulate action from reaction
- Python's first-class functions makes this easy

## UI



```
$("#action-btn").on("click", ...)
```

## Django

auth.user\_logged\_in etc.

models

## SQLAlchemy

- session events
- attribute/model updates

## Tradeoffs

- decouple action from reaction: flexibility & focus
- widely applicable
- can lead to complex interactions & difficult to trace cascades

\* \* \*

Postgres does this!

Look at how Postgres implements Observer

## PostgreSQL: NOTIFY/LISTEN

DB connected clients can

- LISTEN on a message channel
- NOTIFY other listeners of events
- not directly tied to DB objects (relations, tuples, etc.)

## Channeling


LISTEN anything\_good\_on\_this

- channel name is an *identifier*
- 63 char effective limit (w/o recompilation), truncates
- quoted names are case sensitive, unquoted names are not (!)

## “Message for you, sir!”

NOTIFY arthur\_king\_of\_the\_britons, 'ARROWED!'

## psql NOTIFY/LISTEN



## Transaction Interaction

```
BEGIN;
NOTIFY unspeakable_thoughts, 'I'm reluctant to tell you this...'
ROLLBACK;
```

## Transaction Interaction

```
LISTEN im_sorry_you_said_something;
```

```
BEGIN;
```

```
-- NOTIFY posted here
```

```
SELECT some_data FROM this_monster_view;
UPDATE table_with_a_zillion_indexes...;
```

```
COMMIT;
```

```
-- message delivered here
```

## Typical...


```
CREATE FUNCTION new_user_handler() RETURNS TRIGGER AS $$
BEGIN
    EXECUTE 'NOTIFY new_user;';
    RETURN NULL;
END;
$$ LANGUAGE plpgsql;
```

```
CREATE TRIGGER app_user_trigger
AFTER INSERT ON app_user
FOR EACH ROW EXECUTE PROCEDURE new_user_handler();
```

## Use Case

- async job processing
- separate from app server
- e.g. reports, batch/cleanup jobs, notifications

## Goals



- keep UI responsive for tasks initiated by user
- minimize perceived impact on system

## Why [ ! \$OTHER\_TECH ]

### MEANWHILE IN THE IT DEPARTMENT



DevOps has work to do. Duh.

Leverage existing infrastructure

## Data Proximity

Dovetail with TRIGGER

## Approach

- user action/automated tasks create jobs
- job triggers notice
- worker farm picks up new job & runs it

## Trigger Happy

```
CREATE FUNCTION job_queued_handler() RETURNS TRIGGER AS $$
BEGIN
    EXECUTE 'NOTIFY job_queued, '' ' || NEW.queue_name || ''';
    RETURN NULL;
END;
```

```
END;
$$ LANGUAGE plpgsql;

CREATE TRIGGER job_queued_trigger
AFTER INSERT ON job_queue
FOR EACH ROW EXECUTE PROCEDURE job_queued_handler();
```

NEW is the row that has been inserted

## PyListen

```
conn = psycopg2.connect(DSN)
conn.set_isolation_level(ISOLATION_LEVEL_AUTOCOMMIT)
cursor = conn.cursor()

cursor.execute('LISTEN job_queue_control_message;')
cursor.execute('LISTEN job_queued;')

while self.running:
    if select.select([conn], [], [], 10) != ([], [], []):
        conn.poll()
        while conn.notifies:
            notice = conn.notifies.pop()
            print("Got NOTIFY from PID {}: {} {}".format(notice.pid, notice.channel, notice.payload))
```

## Get the Word Out

```
def dblisten(q):
    # connect/loop as above
    q.put(notice)

@websocket.WebSocketWSGI
def broadcast_notify(ws):
    "Relay PG NOTIFY to WebSocket"

    q = eventlet.Queue()
    eventlet.spawn(dblisten, q)
    while True:
        notice = q.get()
        ws.send(unicode(notice.payload))
```

or gevent, ws4py, ØMQ, Redis PUB/SUB, etc.

## Ideas

- notify active subscribers
- streaming monitoring/analytics
- cache invalidation
- µservice architecture
- notify active user/post subscribers
- streaming analytics
- monitoring/event notification
- cache invalidation
- µservice architecture
- biz KPIs

## Performance Notes

- in production for 1.5+ years, millions of jobs
- many concurrent notices  $\rightarrow n^2$  notify performance
  - mitigate w/ relay



@drocco007